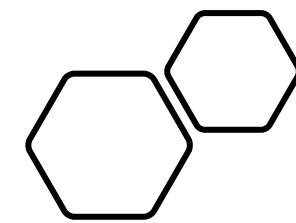
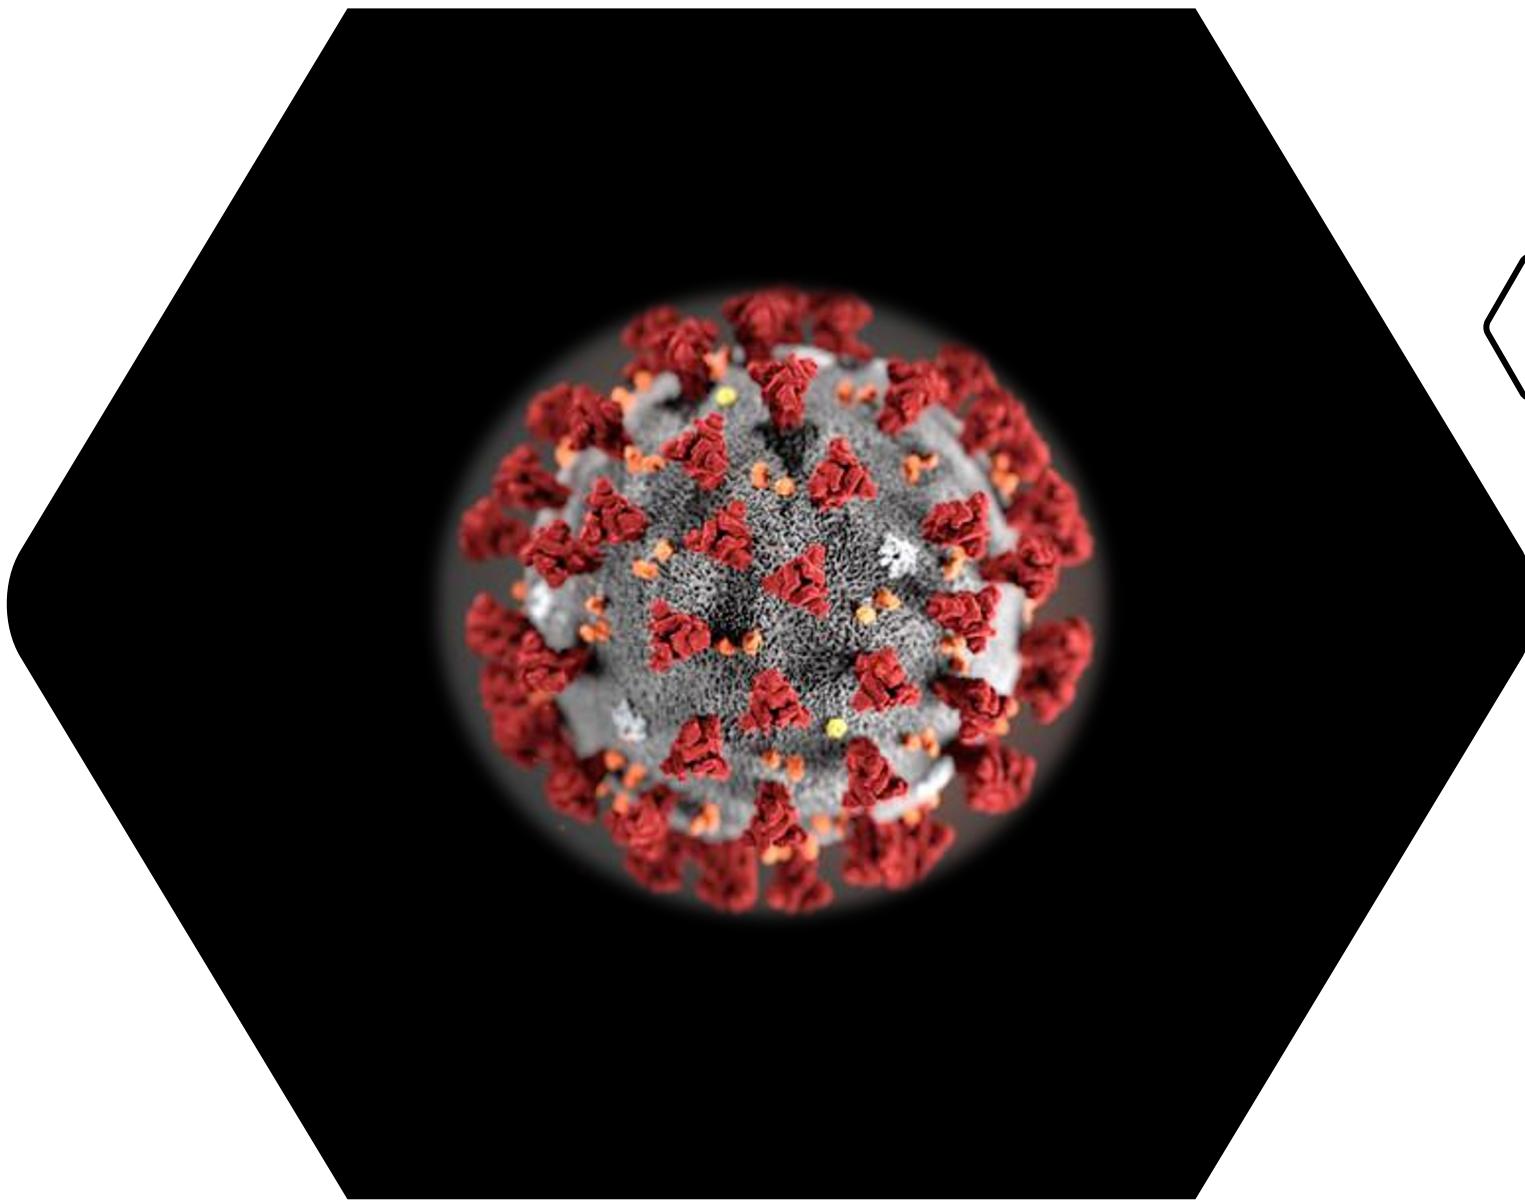




Beyond CS106A

Chris Piech
CS106A, Stanford University

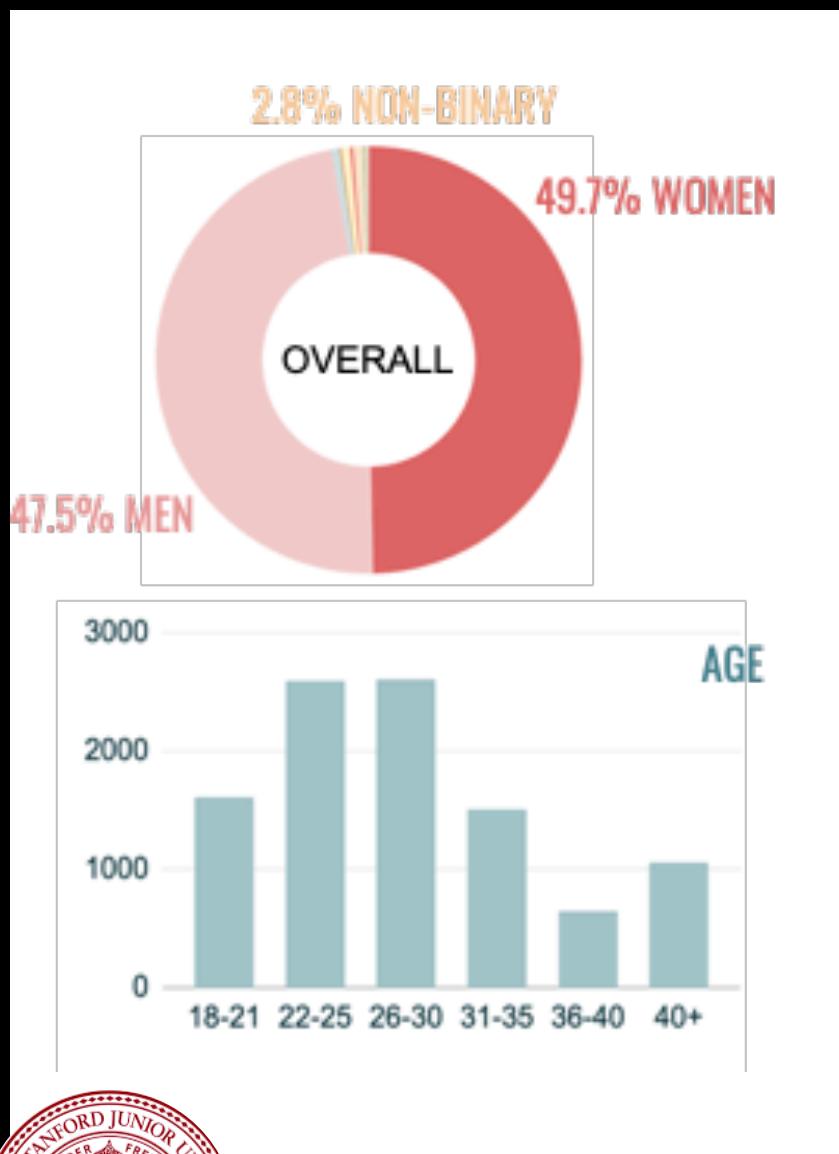
Did you know?



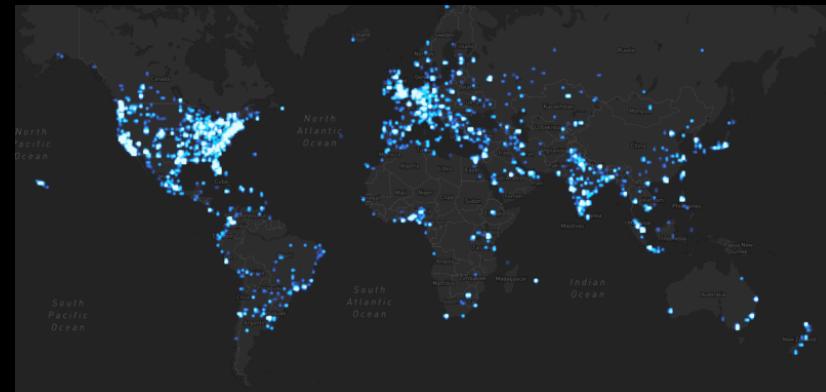
Piech and Sahami, Code in Place



Code in Place : Course with the most section leaders?



**905 section leaders teach
10,000 students
First half of Stanford CS106A**



**20% experienced job loss or home loss
10x retention vs baseline MOOC
99% wanted more (after first section)
6k hours of live teaching
60k hours of lecture watched
Better CS106A for Stanford students**

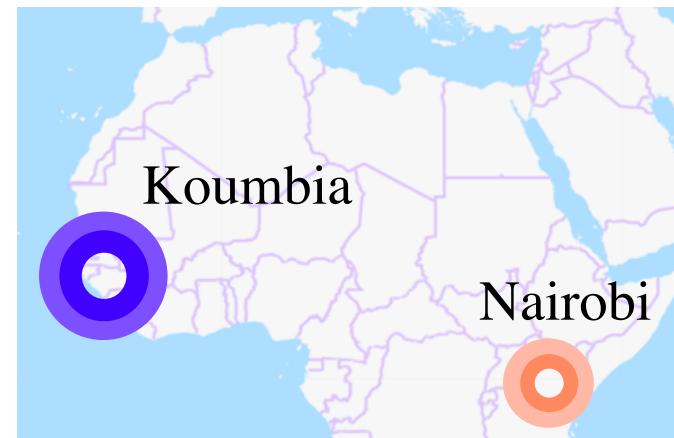


Have you thought about teaching?

Great way to learn

Great way to give back

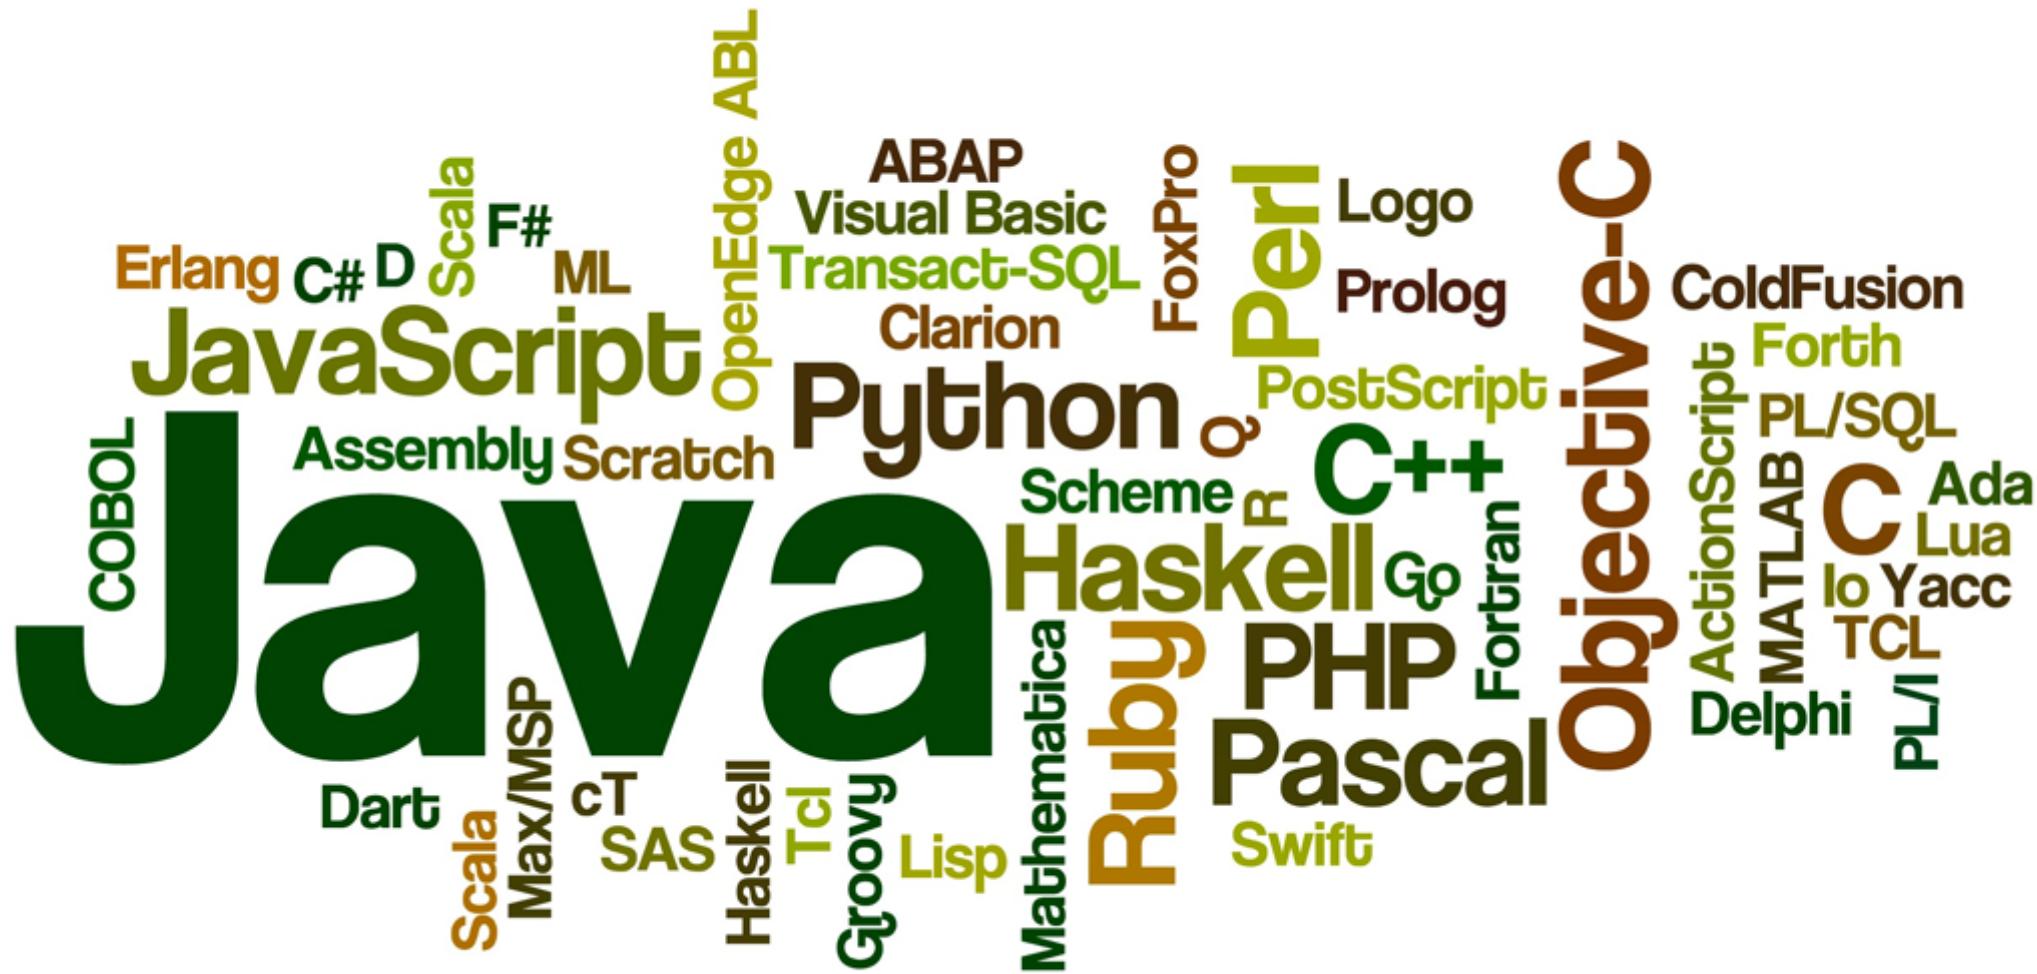
In-person teaching around the world.



1. How to make your own project
2. What other languages look like
3. Deep Learning in Python

Life after CS106A

Programming Languages



Python

```
evens = []
for i in range(100):
    if i % 2 == 0:
        evens.append(i)
print evens
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



C++

```
Vector<double> evens;
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
cout << evens << endl;
```

prints [2, 4, 6, 8, 10, 12, ...]



Java

```
ArrayList<Double> evens = new ArrayList<Double>();
for(int i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.add(i);
    }
}
println(evens);
```

prints [2, 4, 6, 8, 10, 12, ...]



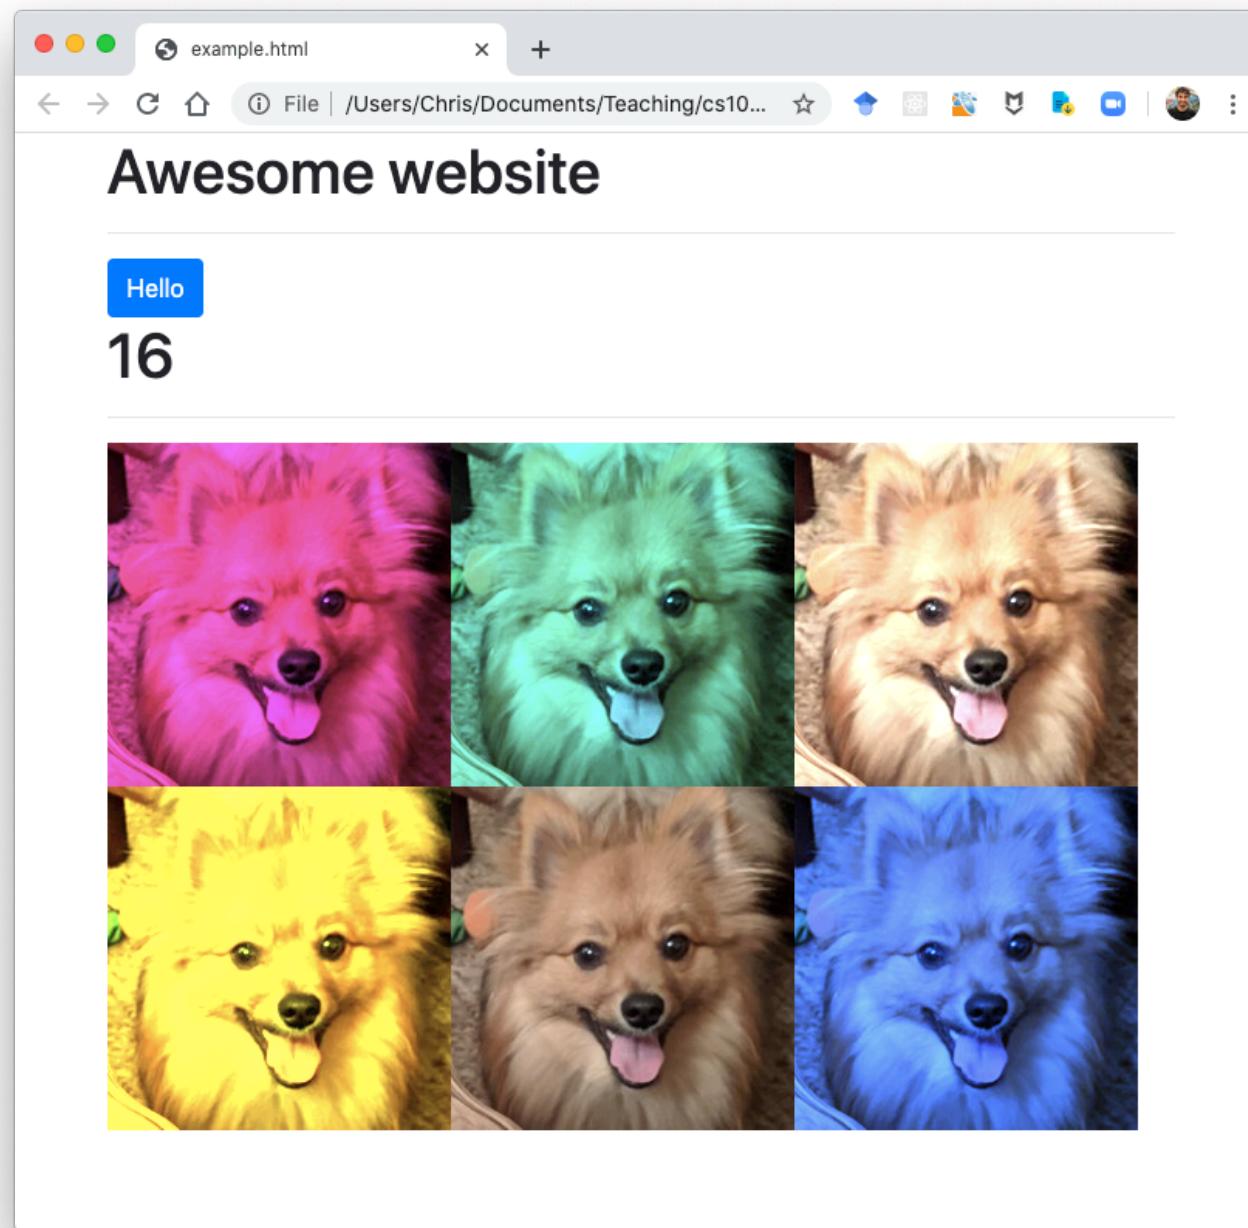
Javascript

```
var evens = []
for(var i = 0; i < 100; i++) {
    if(i % 2 == 0) {
        evens.push(i)
    }
}
console.log(evens)
```

prints [2, 4, 6, 8, 10, 12, ...]



Lets Play



There is something going on
in the world of AI

[suspense]

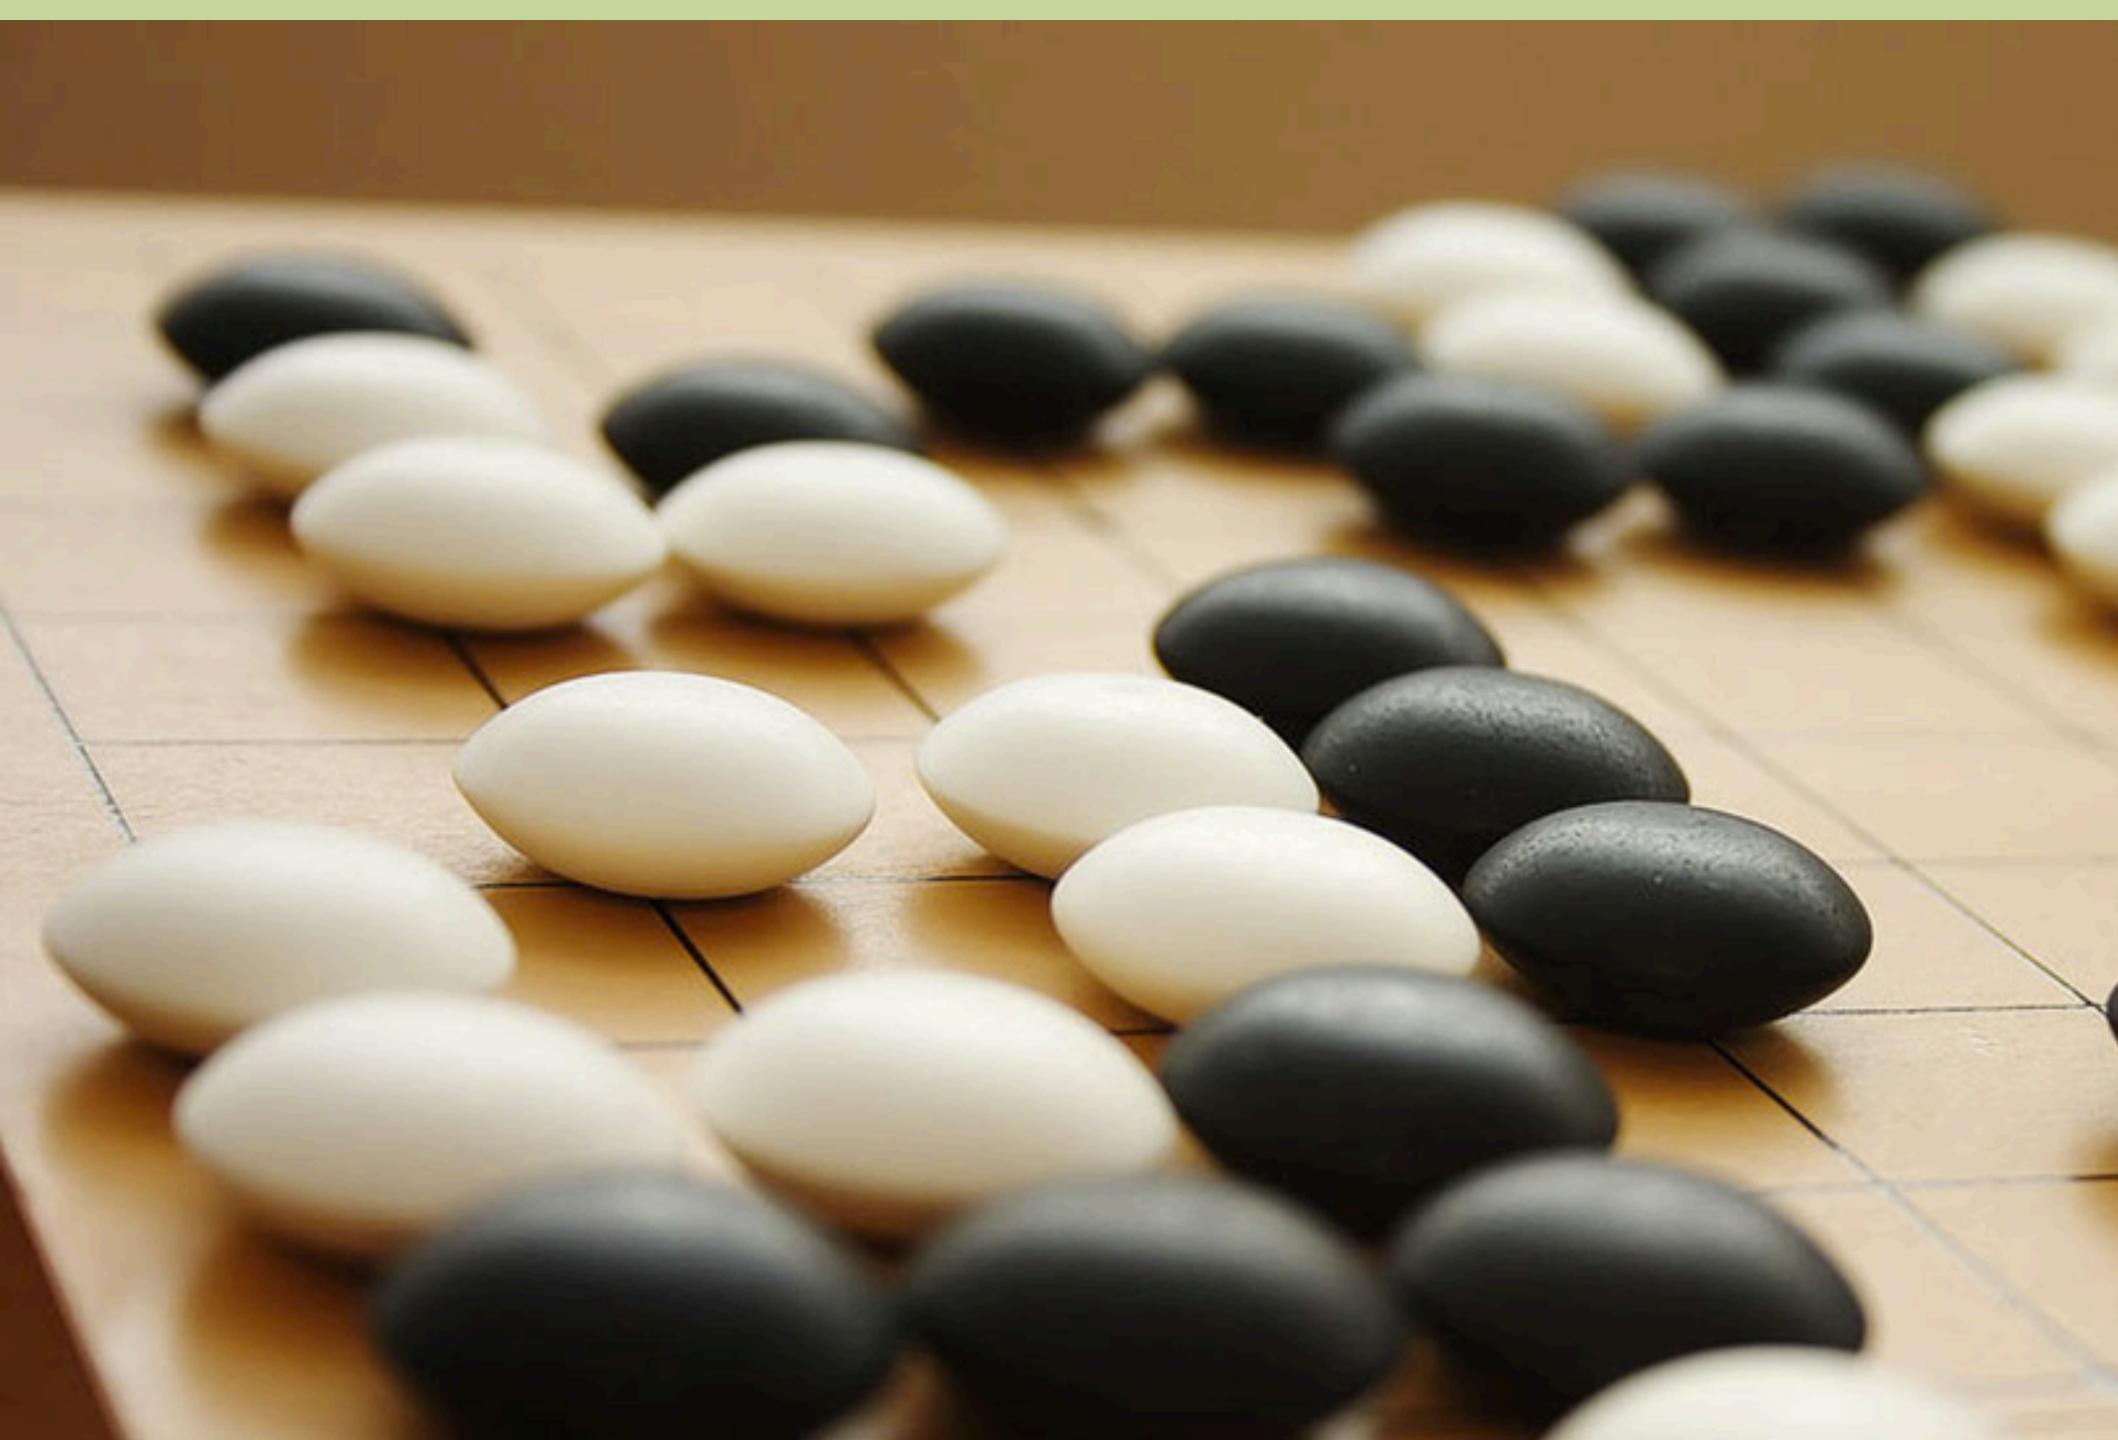
Self Driving Cars



Computers Making Art



The Last Remaining Board Game

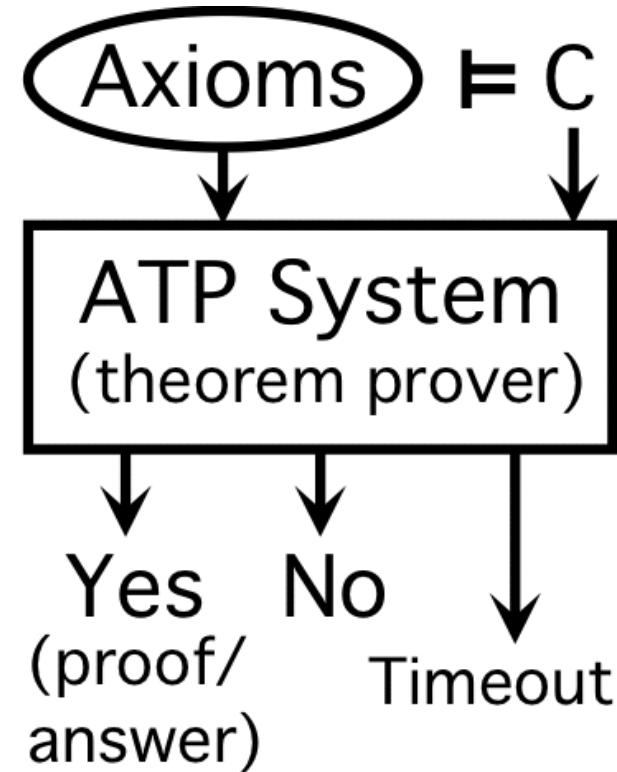


Early Optimism 1950

1952



1955



Computer Vision



Piech, CS106A, Stanford University



Classification



Classification



Classification



* It doesn't have to
be correct all of the
time



Identifying Cats

Here's one way you might code this...

```
def is_cat(image):
    if contains_two_eyes(image):
        if has_whiskers(image):
            if has_pointy_ears(image):
                return True
    return False
```



Some Tricky Cases

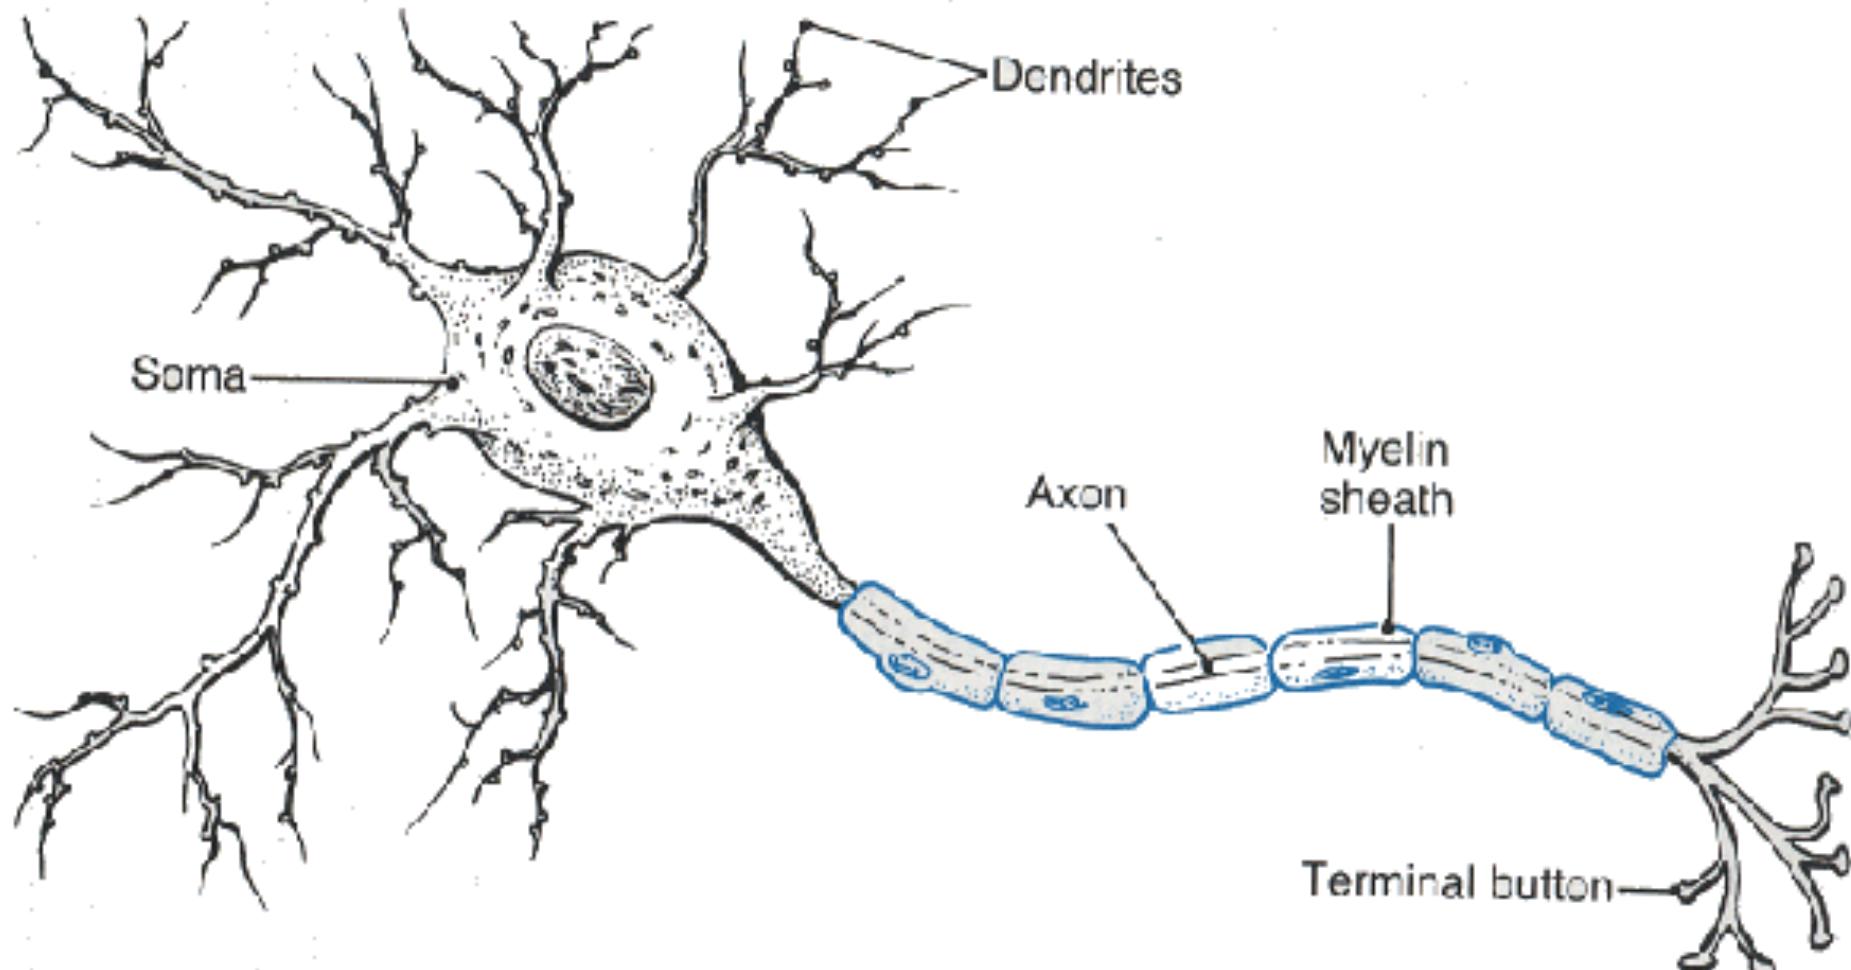


Meen, CS100A, Stanford University

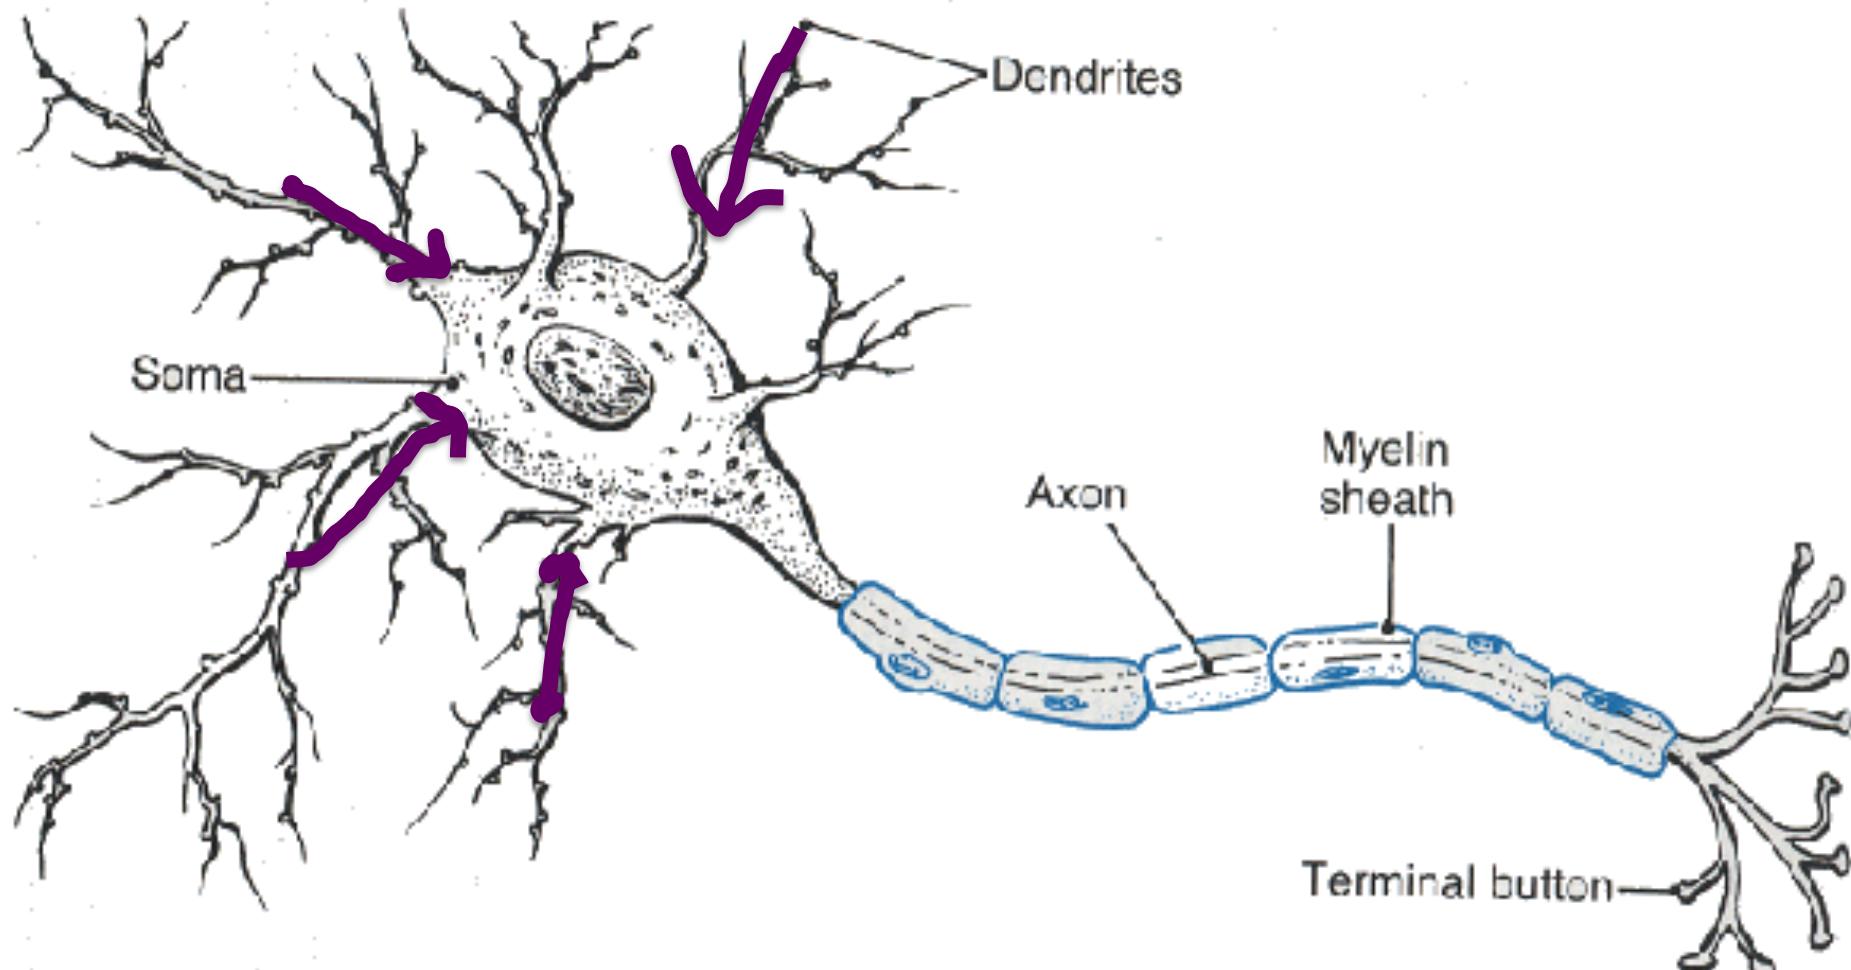


Great idea inspired by biology

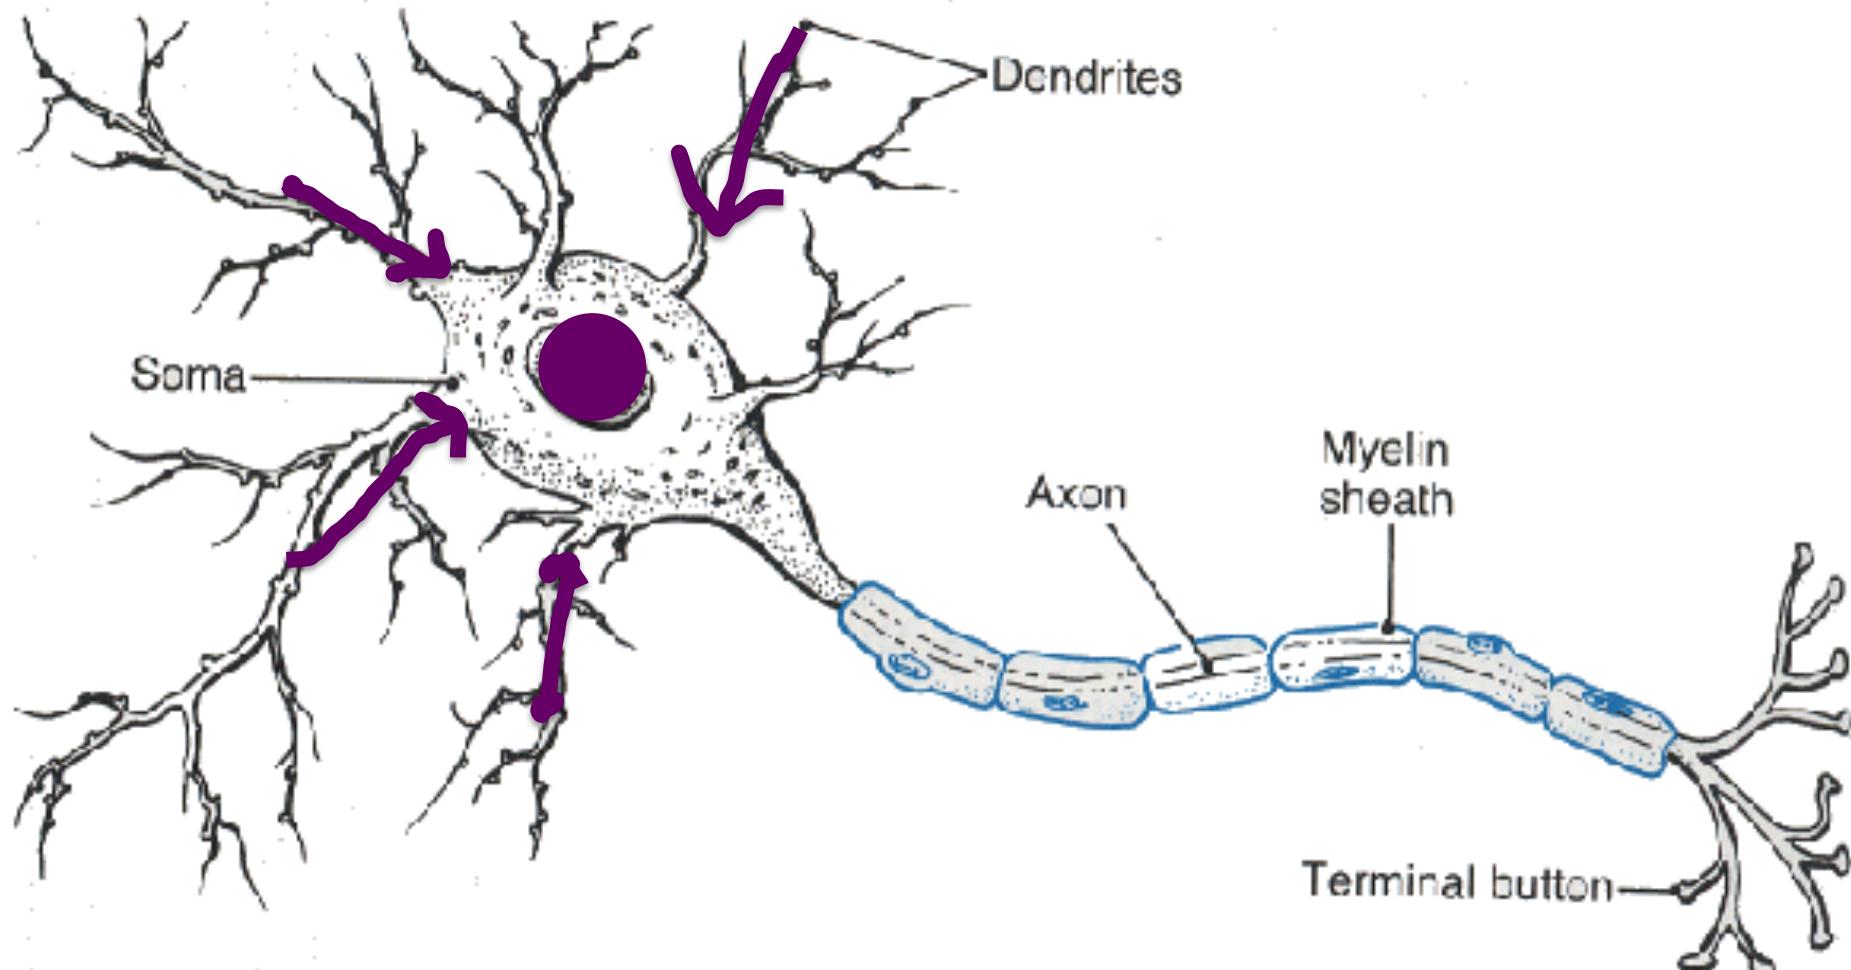
Neuron



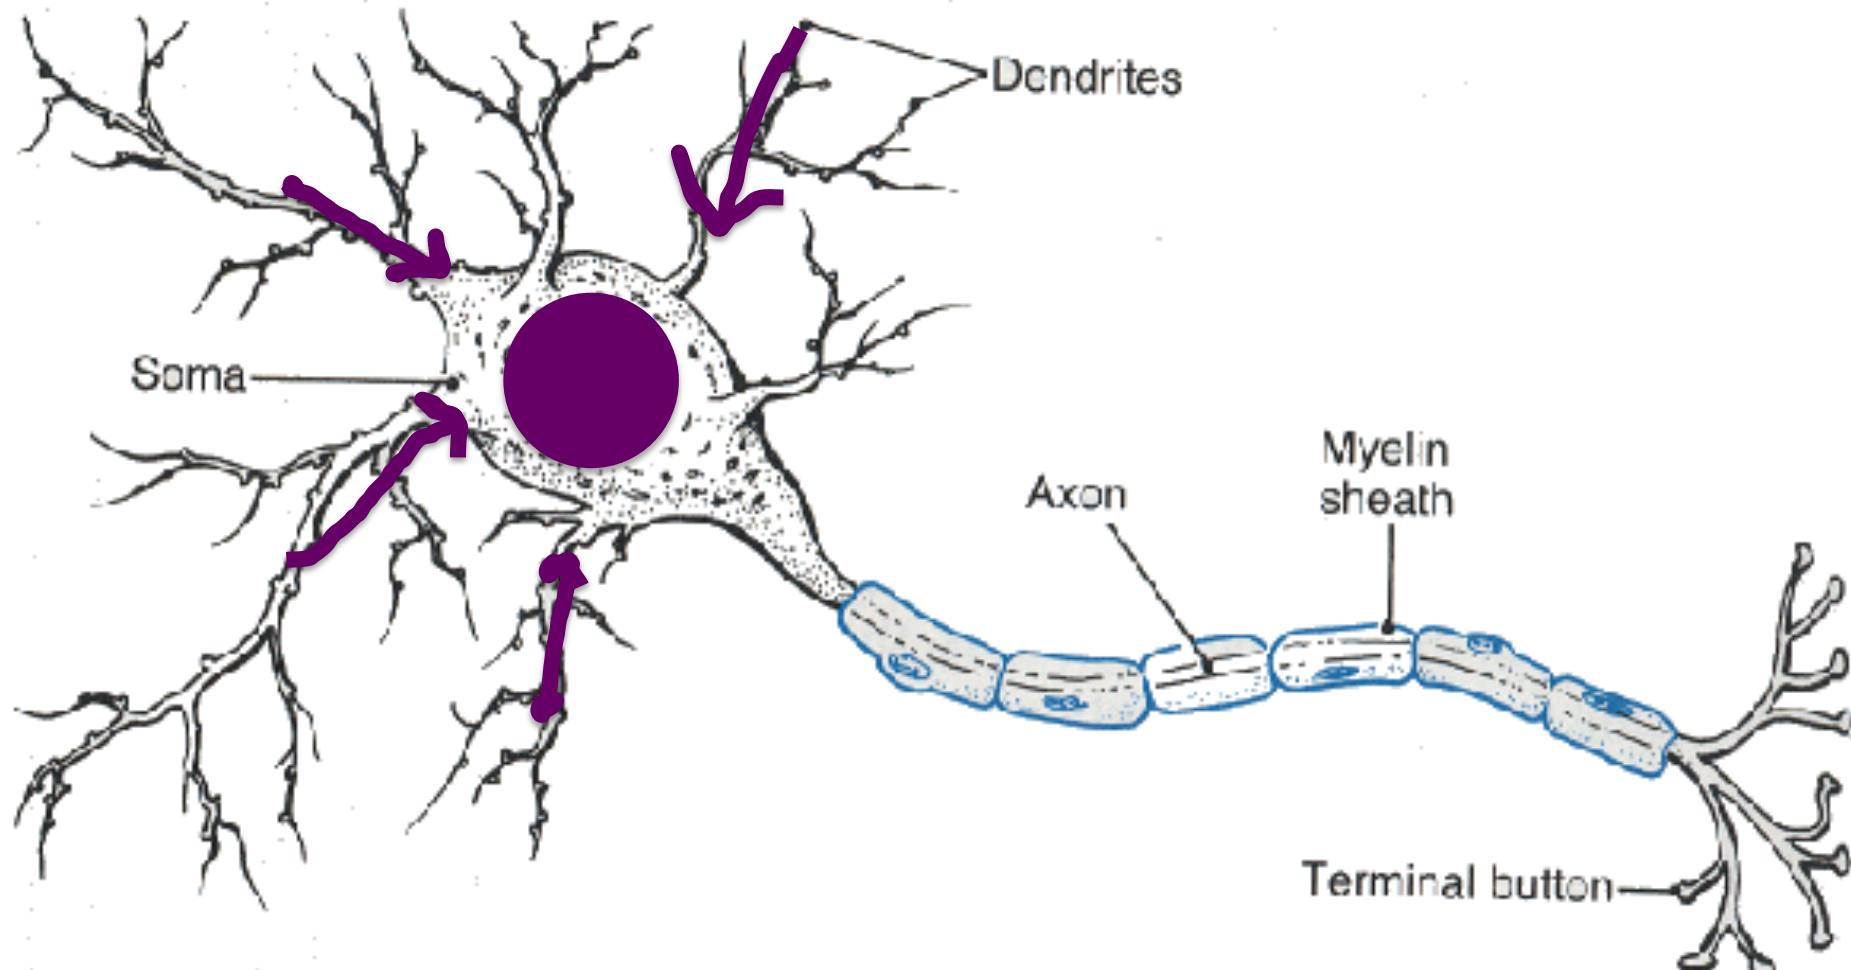
Neuron



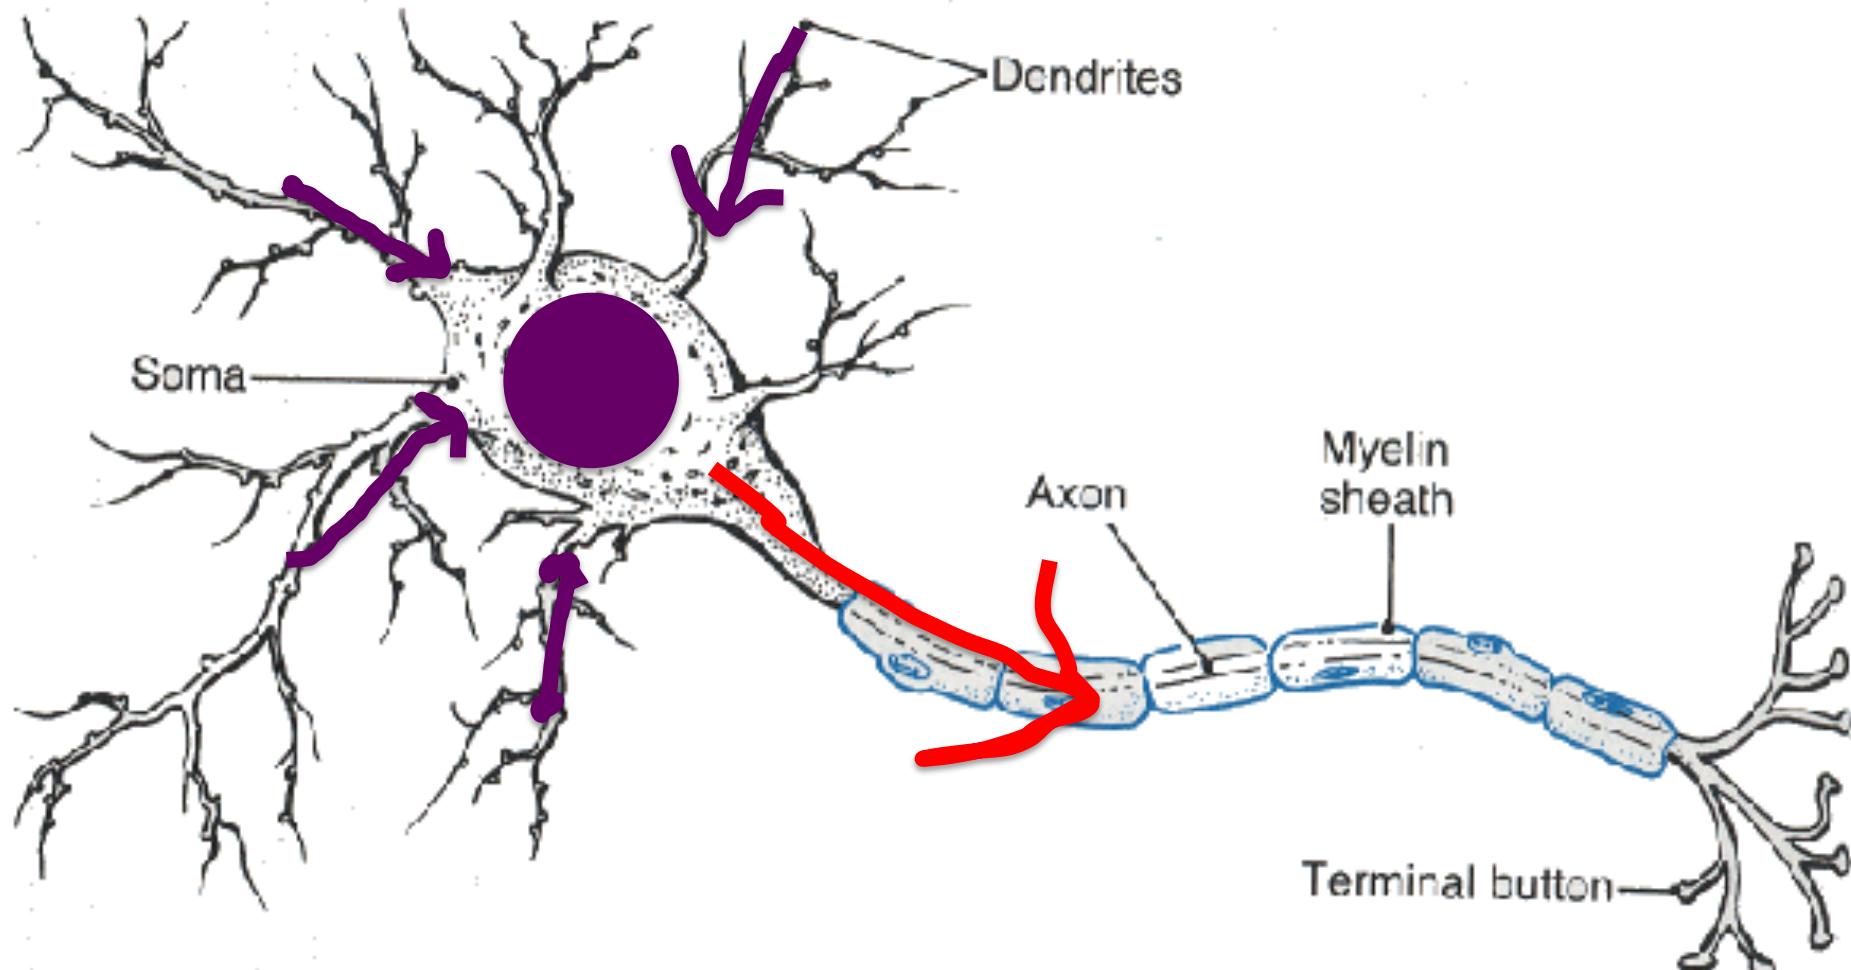
Neuron



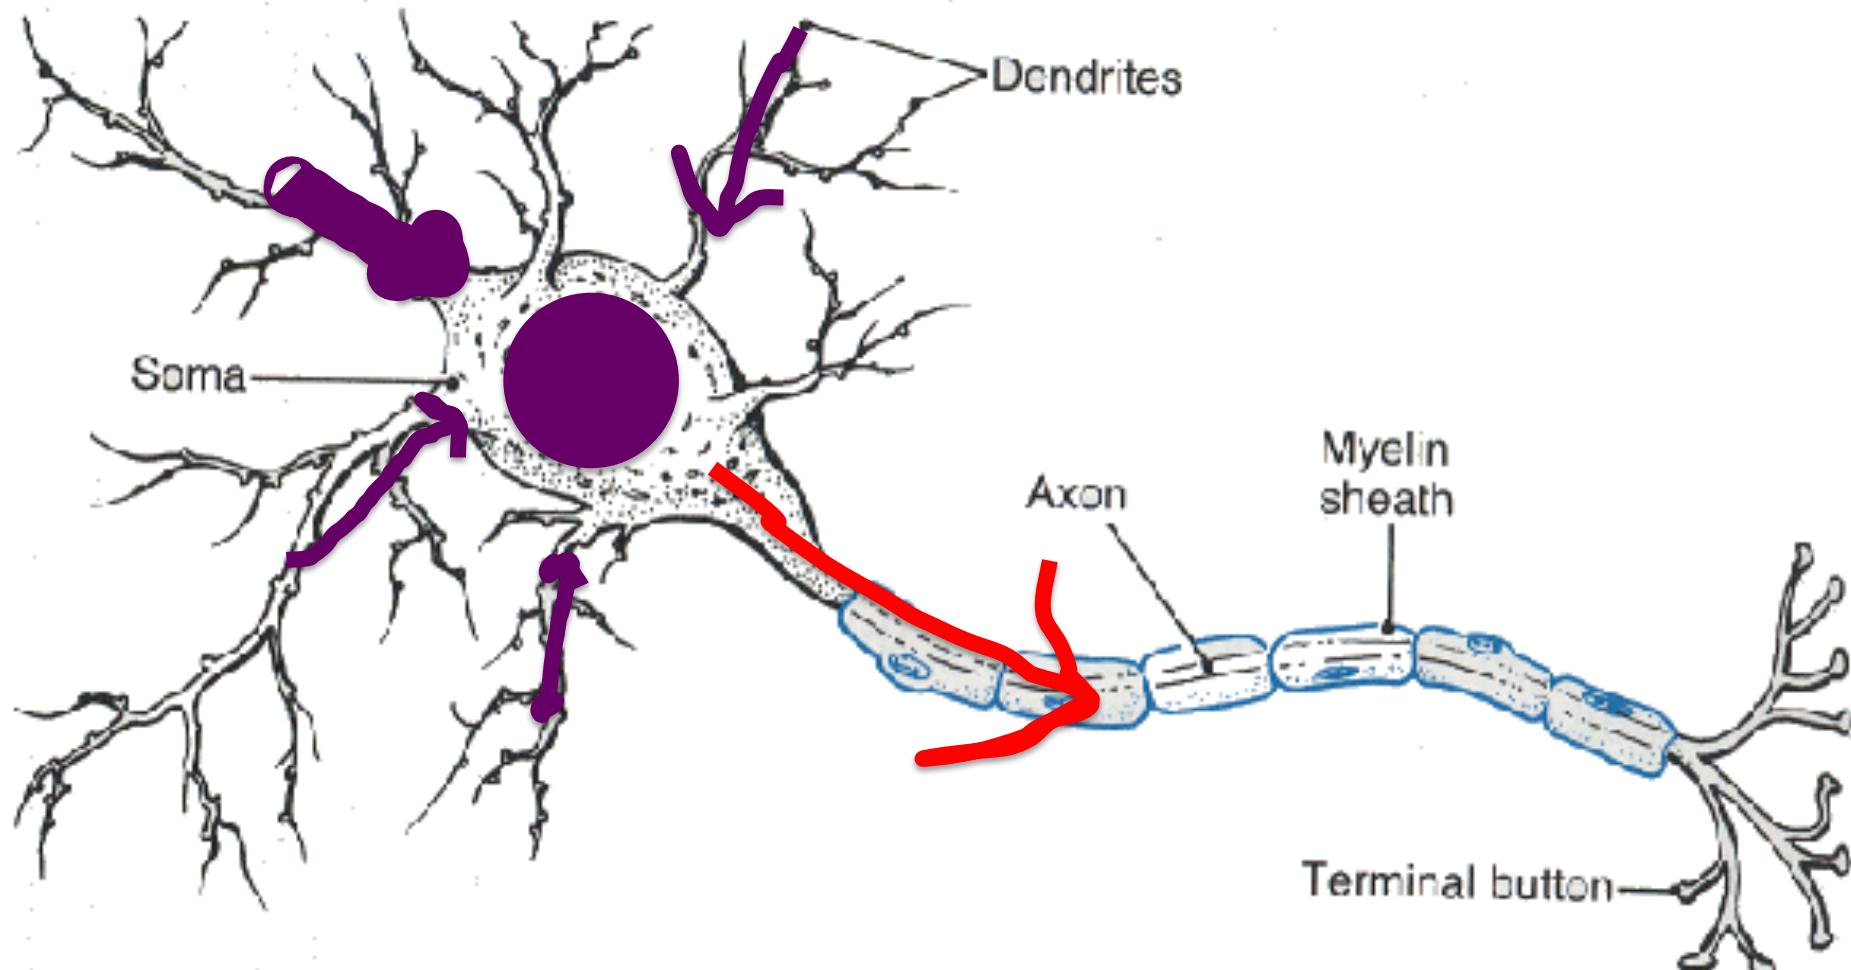
Neuron



Neuron



Some Inputs are More Important

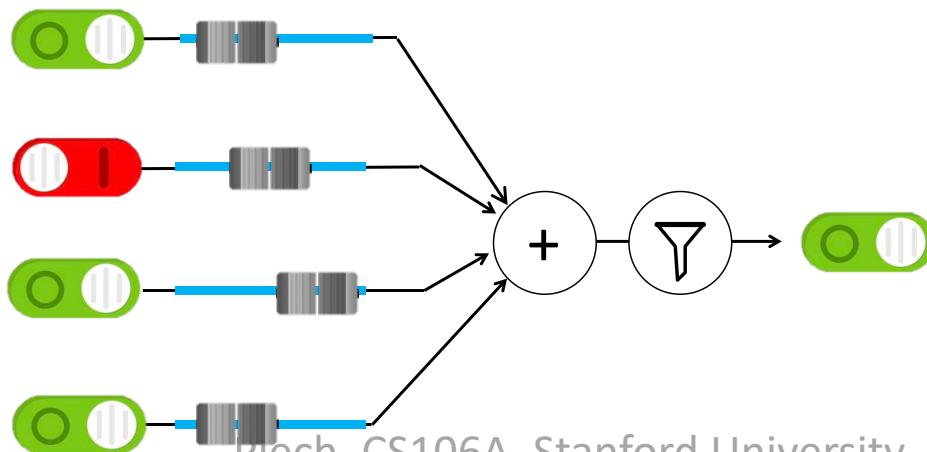


Artificial Neuron

```
# calculate the activation of a neuron
def activate(weights_list, inputs_list):
    weighted_sum = 0;
    for i in range(len(inputs)):
        weighted_sum += weights_list[i] * inputs_list[i]

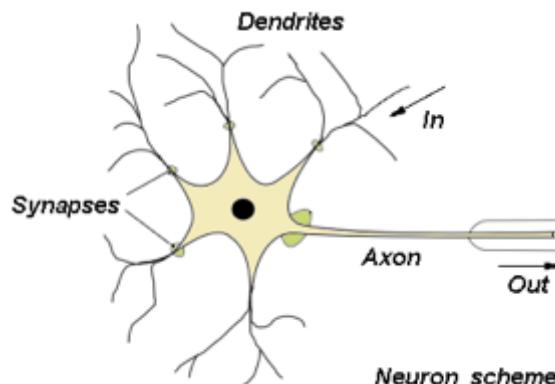
    return squash(weighted_sum);
```

```
# the sigmoid function forces a value to be between 0 and 1
def squash(value):
    return 1 / (1 + math.exp(-value));
```

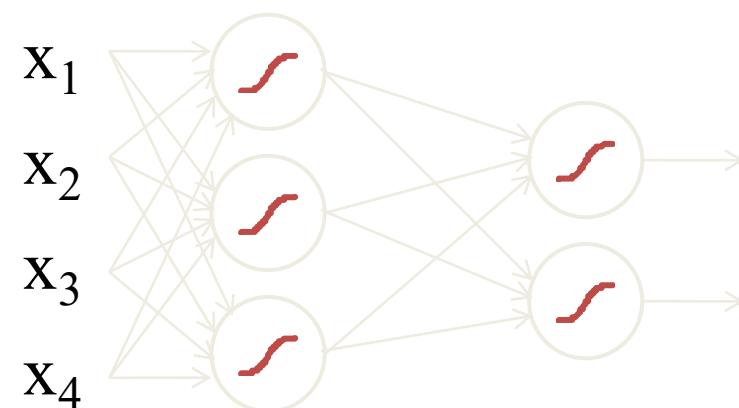
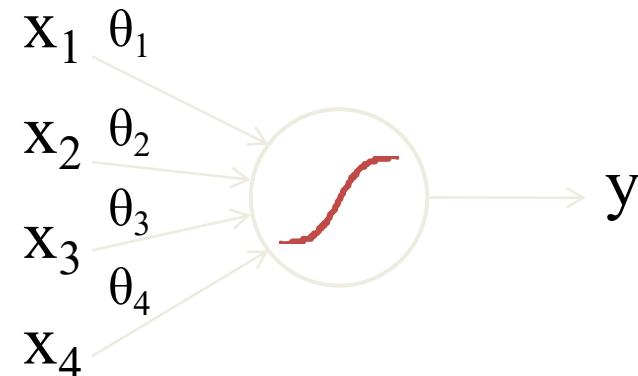
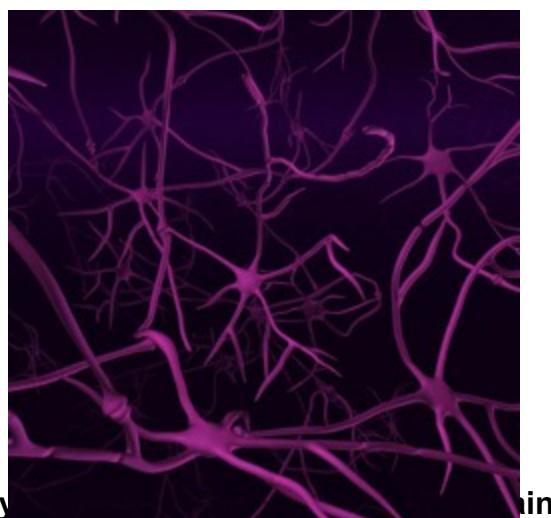


Biological Basis for Neural Networks

- A neuron



- Your brain



Demonstration

Draw your number here



Downsampled drawing:

First guess:

Second guess:

Layer visibility

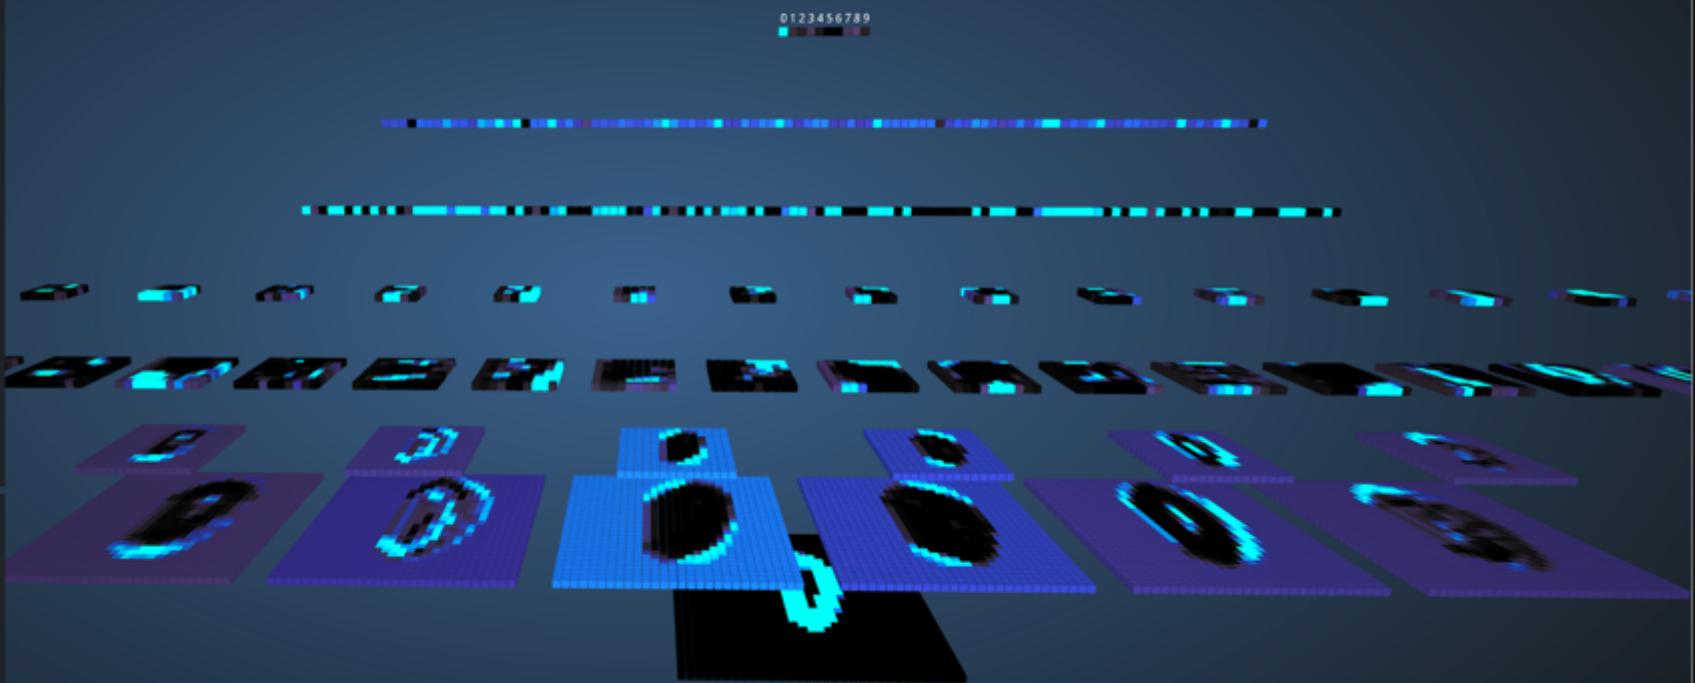
Input layer

Convolution layer 1

Downsampling layer 1

Convolution layer 2

Downsampling layer 2



<http://scs.ryerson.ca/~aharley/vis/conv/>

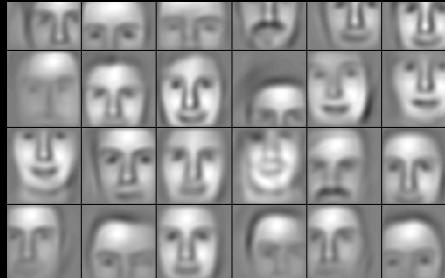
Piech, CS106A, Stanford University



Visualize the Weights



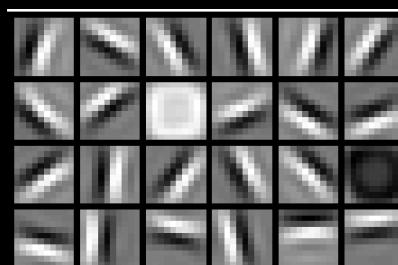
Training set: Aligned
images of faces.



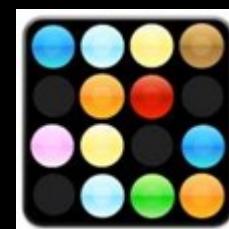
object models



object parts
(combination
of edges)

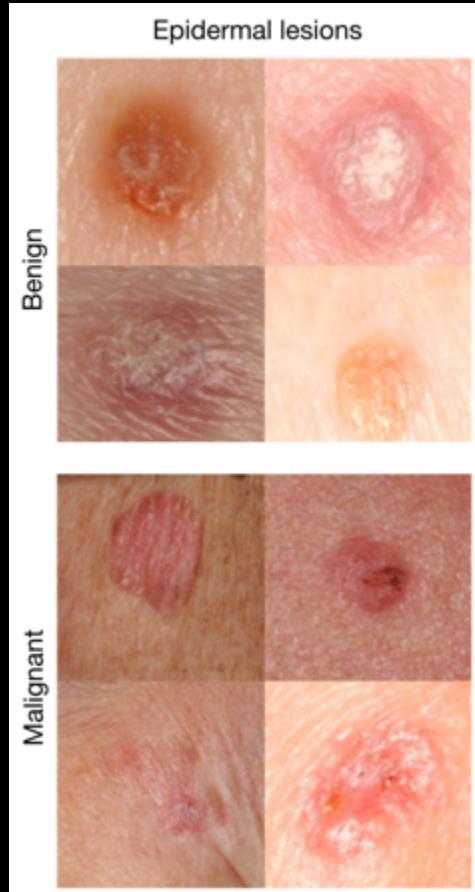


edges



pixels

Where is this useful?

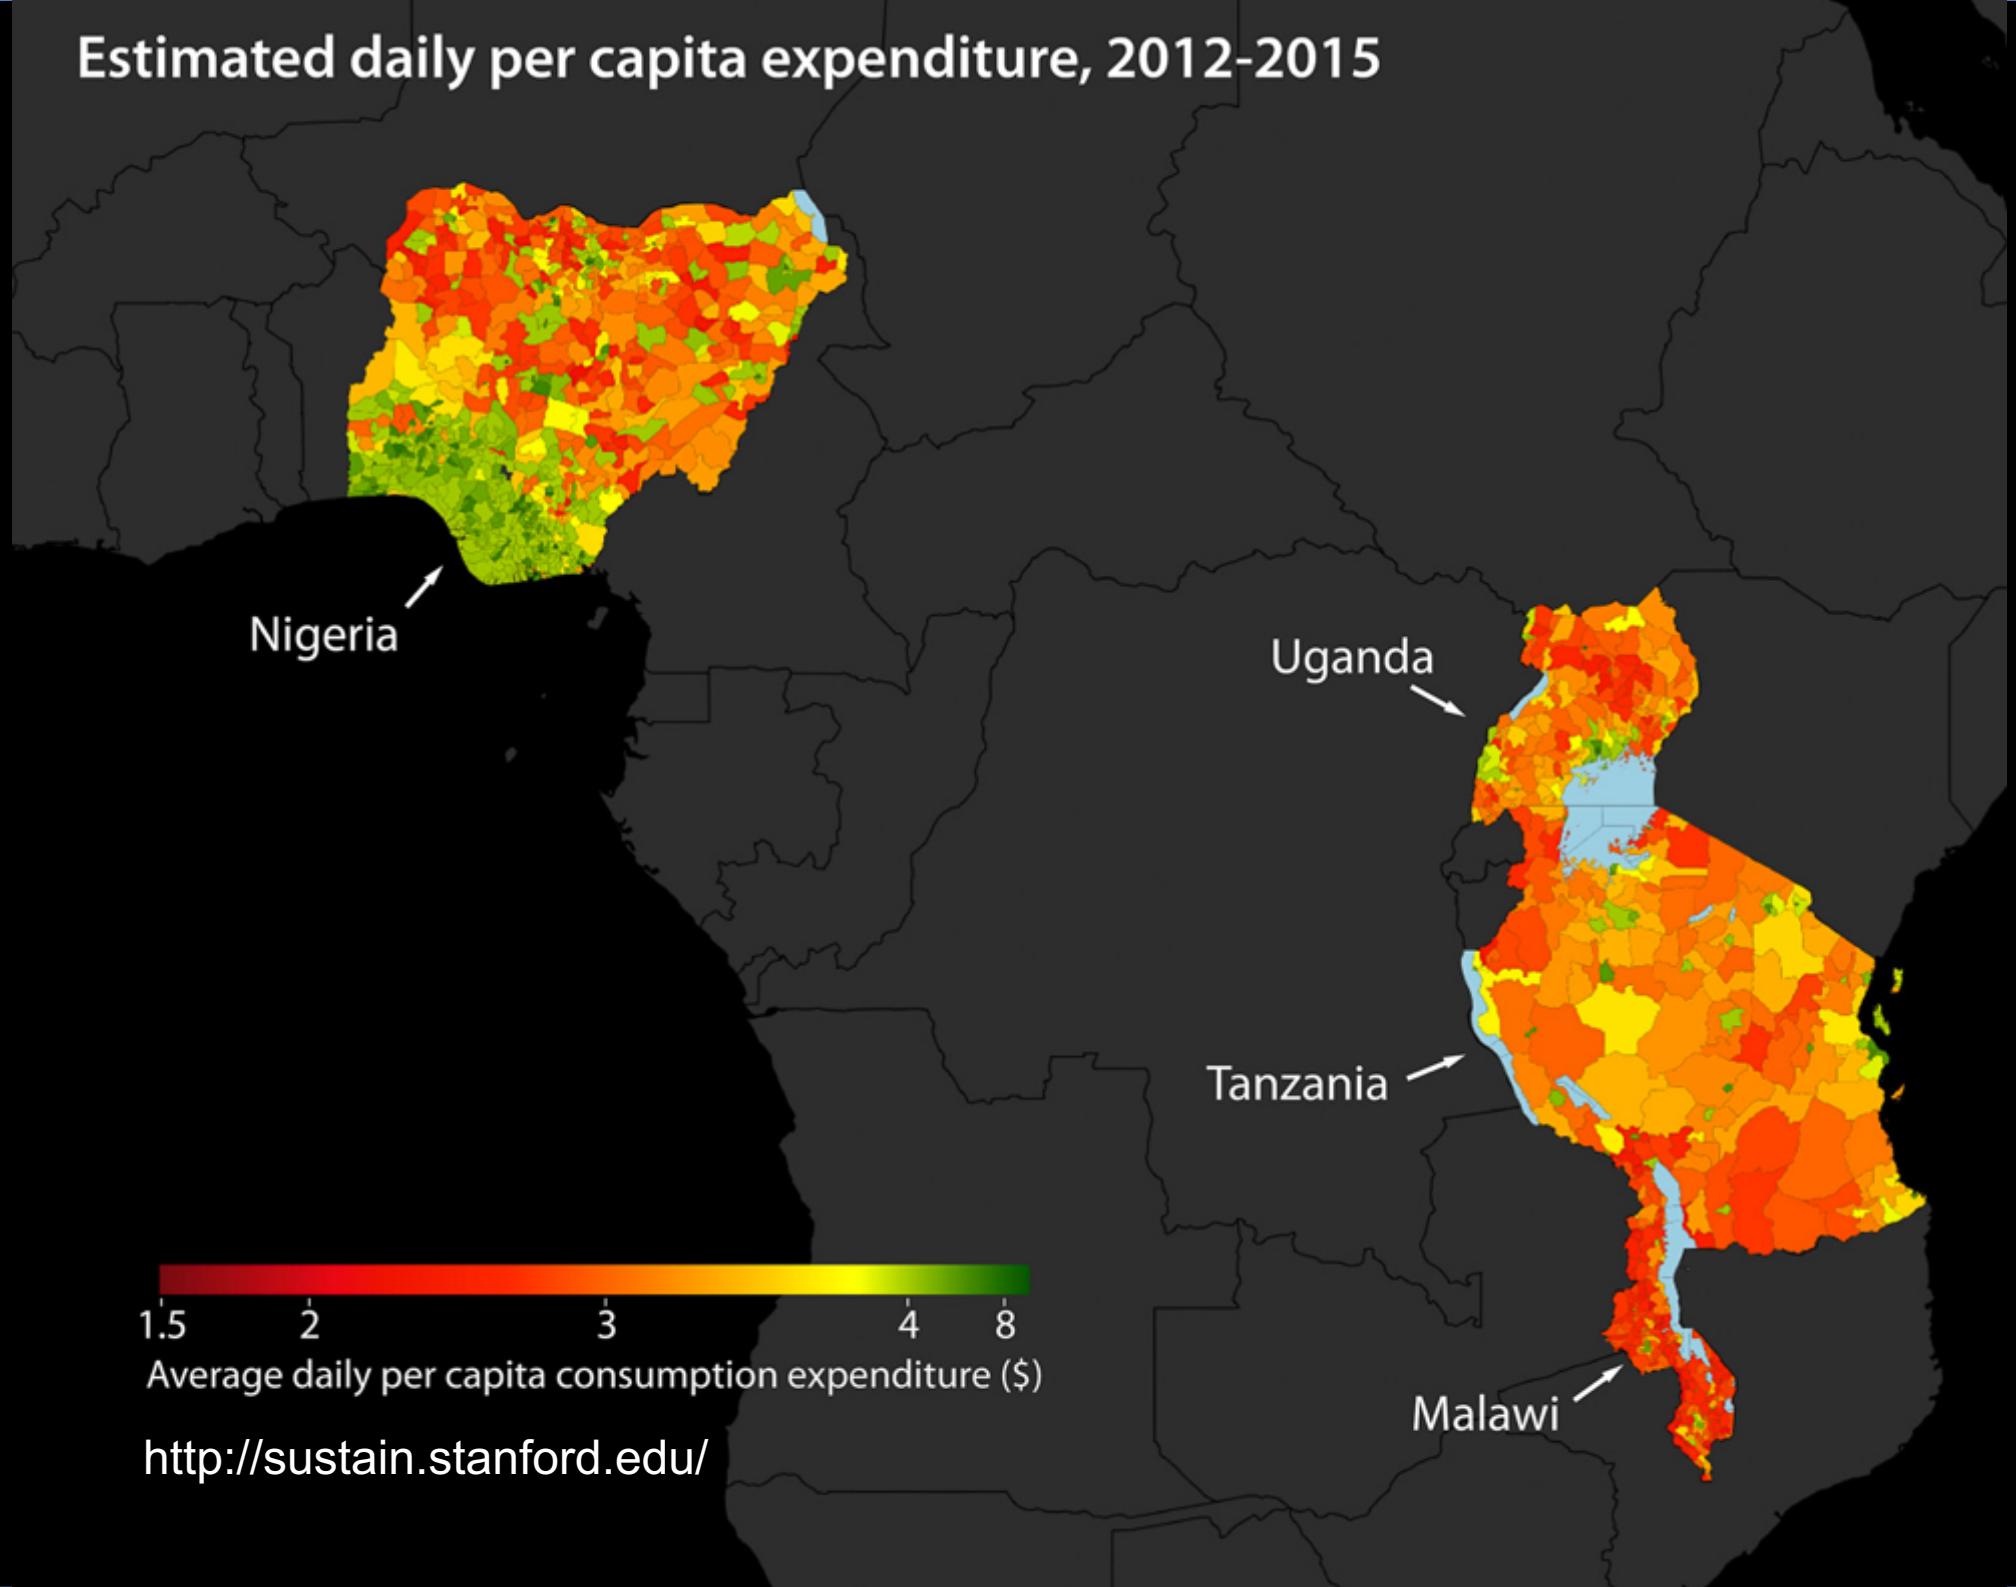


A machine learning algorithm performs **better than** the best dermatologists.

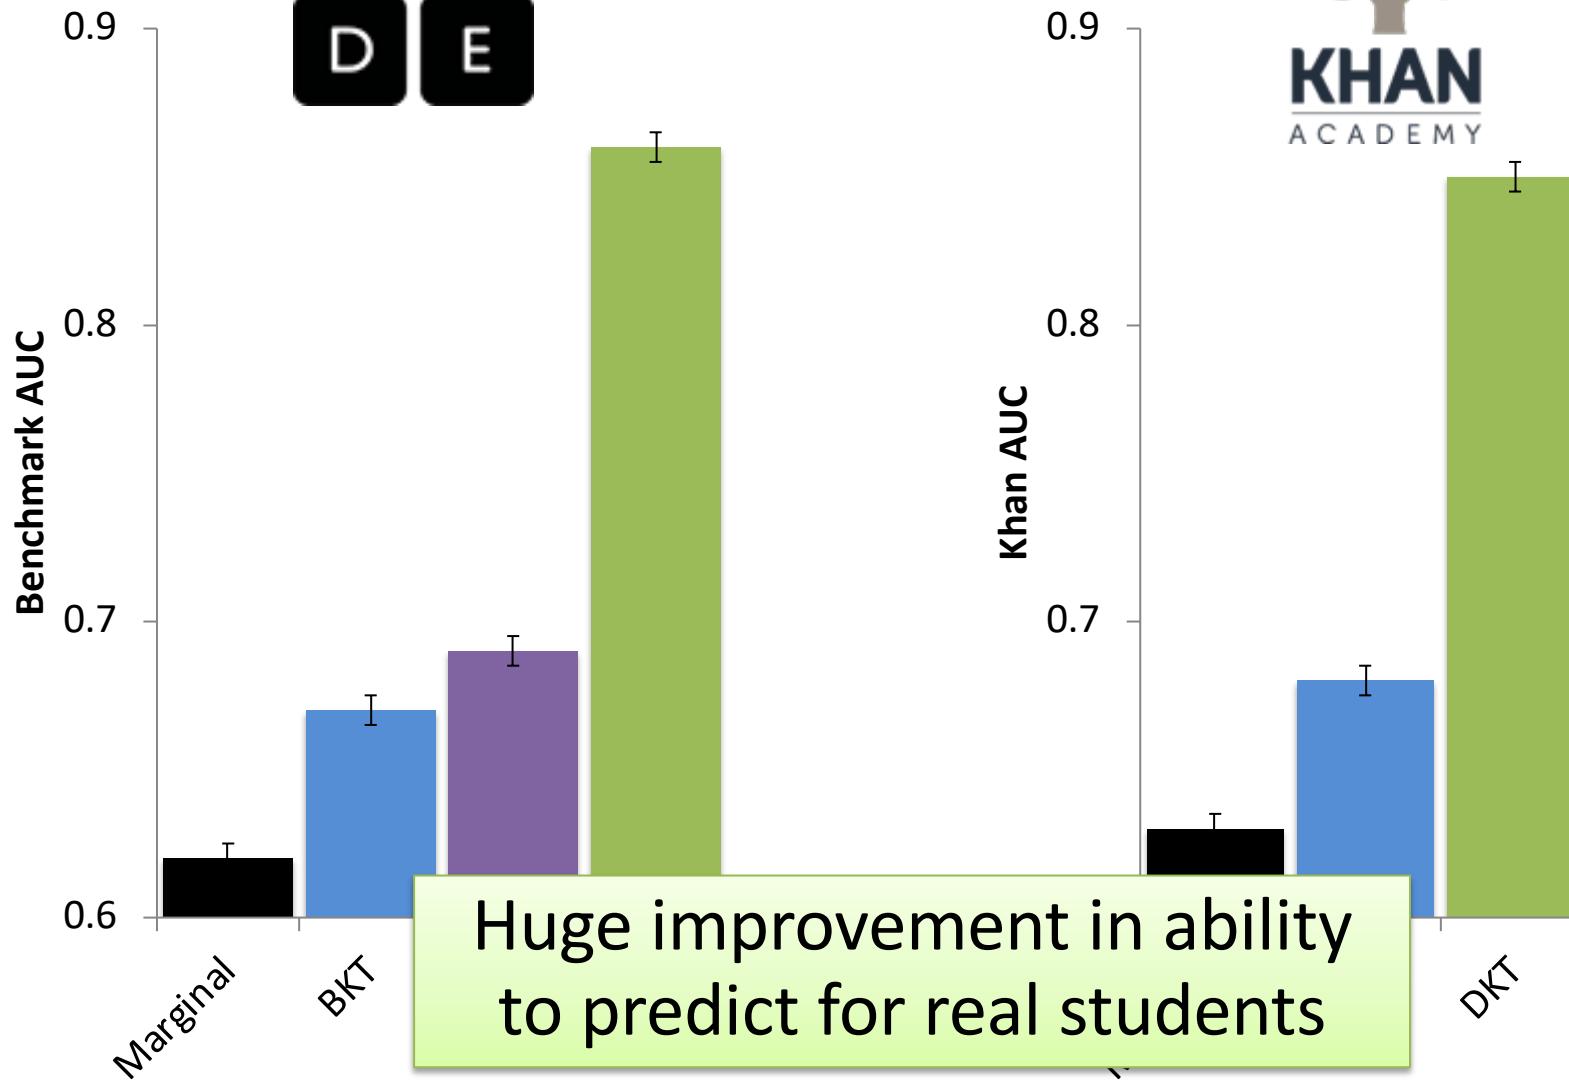
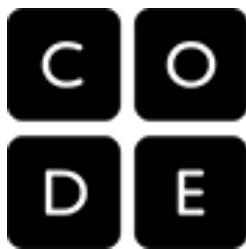
Developed this year, at Stanford.

Esteva, Andre, et al. "Dermatologist-level classification of skin cancer with deep neural networks." *Nature* 542.7639 (2017): 115-118.

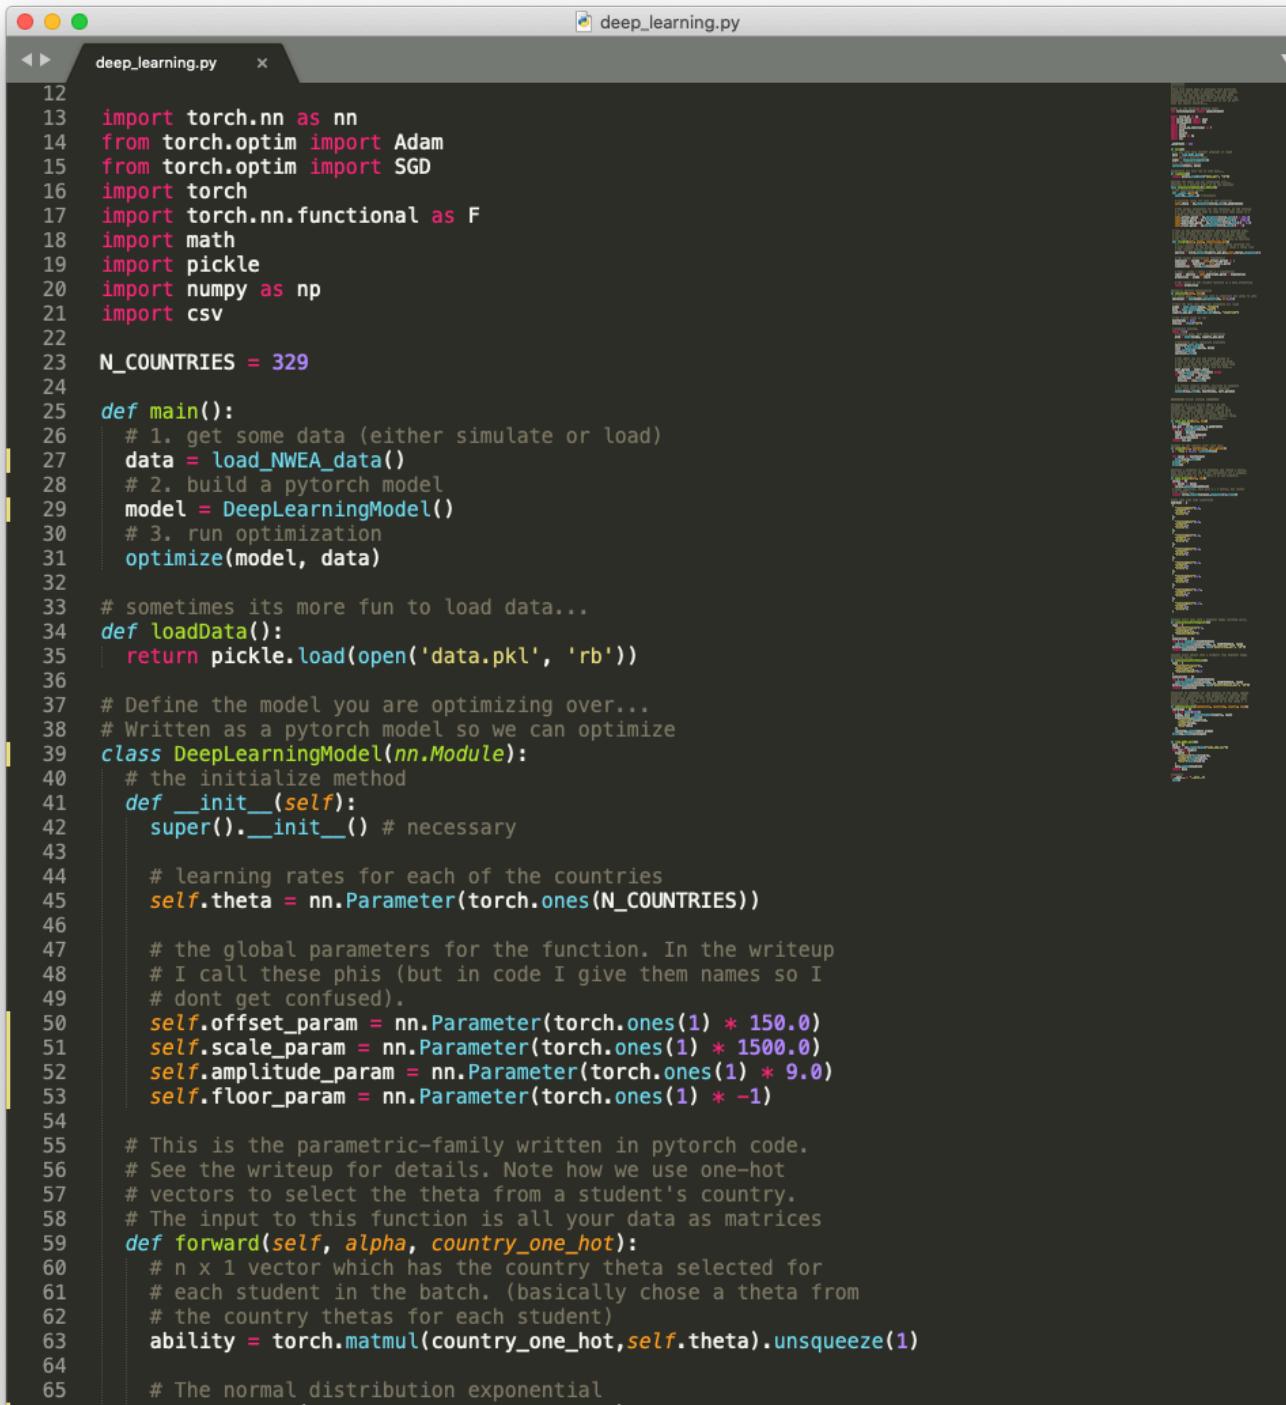
Estimated daily per capita expenditure, 2012-2015



Understanding Students



What does it look like?



A screenshot of a terminal window titled "deep_learning.py". The window displays a large block of Python code. The code imports various PyTorch modules and defines a class named "DeepLearningModel". The "DeepLearningModel" class has several parameters and methods, including a constructor "__init__" and a "forward" method. The "forward" method uses PyTorch's "matmul" function to perform matrix multiplication between "country_one_hot" and "self.theta".

```
12 import torch.nn as nn
13 from torch.optim import Adam
14 from torch.optim import SGD
15 import torch
16 import torch.nn.functional as F
17 import math
18 import pickle
19 import numpy as np
20 import csv
21
22 N_COUNTRIES = 329
23
24 def main():
25     # 1. get some data (either simulate or load)
26     data = loadData()
27     # 2. build a pytorch model
28     model = DeepLearningModel()
29     # 3. run optimization
30     optimize(model, data)
31
32     # sometimes its more fun to load data...
33     def loadData():
34         return pickle.load(open('data.pkl', 'rb'))
35
36     # Define the model you are optimizing over...
37     # Written as a pytorch model so we can optimize
38     class DeepLearningModel(nn.Module):
39         # the initialize method
40         def __init__(self):
41             super().__init__() # necessary
42
43             # learning rates for each of the countries
44             self.theta = nn.Parameter(torch.ones(N_COUNTRIES))
45
46             # the global parameters for the function. In the writeup
47             # I call these phis (but in code I give them names so I
48             # dont get confused).
49             self.offset_param = nn.Parameter(torch.ones(1) * 150.0)
50             self.scale_param = nn.Parameter(torch.ones(1) * 1500.0)
51             self.amplitude_param = nn.Parameter(torch.ones(1) * 9.0)
52             self.floor_param = nn.Parameter(torch.ones(1) * -1)
53
54             # This is the parametric-family written in pytorch code.
55             # See the writeup for details. Note how we use one-hot
56             # vectors to select the theta from a student's country.
57             # The input to this function is all your data as matrices
58             def forward(self, alpha, country_one_hot):
59                 # n x 1 vector which has the country theta selected for
60                 # each student in the batch. (basically chose a theta from
61                 # the country thetas for each student)
62                 ability = torch.matmul(country_one_hot, self.theta).unsqueeze(1)
63
64                 # The normal distribution exponential .
65
```

